

ISAAC LIAO

iliao@andrew.cmu.edu <https://iliao2345.github.io>

EDUCATION

Carnegie Mellon University PhD in Machine Learning. Advisor: Albert Gu GPA: 3.67/4.0	<i>Aug 2024 - May 2030 (Expected)</i>
Massachusetts Institute of Technology Master of Engineering in EECS. Advisor: Max Tegmark GPA: 5.0/5.0	<i>Sep 2023 - May 2024</i>
Massachusetts Institute of Technology Bachelor of Science, Double major in CS and Physics GPA: 5.0/5.0	<i>Sep 2019 - Jun 2023</i>

RESEARCH EXPERIENCE

Albert Gu Group Graduate Researcher (PhD). Advising: Albert Gu	<i>Aug 2024 - Present</i>
<ul style="list-style-type: none">Research on ARC-AGI Without Pretraining: training small models on target puzzles with inference-time-only learning; achieved 20% solve rate on ARC-AGI-1 with a 76k-parameter model.	
Jane Street Machine Learning Research Intern <ul style="list-style-type: none">Built generative models to infer unseen activity in partially observed markets.Developed ensembling techniques for neural networks trained on market data.	<i>May 2025 - Aug 2025</i>
Tegmark AI Safety Group Graduate Researcher (MEng). Advising: Max Tegmark	<i>Sep 2023 - May 2024</i>
<ul style="list-style-type: none">Discovered circular representations of days of the week in Mistral 7B and GPT-2-small.Simplified recurrent neural networks into standard forms using symmetry transformations.	

PUBLICATIONS

ARC-AGI Without Pretraining <i>Isaac Liao, Albert Gu.</i>	Blog post , also submitted to NeurIPS 2025.
<ul style="list-style-type: none">Solving IQ-test-like visual puzzles by <i>extreme generalization</i>, training only on the target puzzle.20% solve rate on ARC-AGI-1 despite training only at inference time, with a tiny 76K param model.	
Not All Language Model Features Are One-Dimensionally Linear. <i>Josh Engels, Eric J. Michaud, Isaac Liao, Wes Gurnee, and Max Tegmark.</i>	arXiv 2024.
<ul style="list-style-type: none">Discovering circular representations of days of the week in Mistral 7B and GPT-2-small.	
Opening the AI Black Box: Program Synthesis via Mechanistic Interpretability. <i>Eric J. Michaud, Isaac Liao, Vedang Lad, Ziming Liu, Anish Mudide, et al.</i>	arXiv 2024.
<ul style="list-style-type: none">Reducing RNN weights into interpretable python code through a series of simplifying steps.	
Learning to Optimize Quasi-Newton Methods. <i>Isaac Liao, Rumen Dangovski, Jakob Nicolaus Foerster, and Marin Soljačić.</i>	TMLR 2023.
<ul style="list-style-type: none">Online learning a preconditioner for gradient descent; theoretical guarantees on preconditioner behavior.	
Streamlining Physics Problem Generation to Support Physics Teachers in Using Gen. AI. <i>Shams El-Adawy, Isaac Liao, Vedang Lad, Mohamed Abdelhafez, et al.</i>	The Physics Teacher 2024.
<ul style="list-style-type: none">Techniques for chain-of-thought prompting an LLM to generate physics problems suitable for teaching.	
Generating Interpretable Networks Using Hypernetworks. <i>Isaac Liao, Ziming Liu, and Max Tegmark.</i>	arXiv 2023.
<ul style="list-style-type: none">Designing a graph neural network to generate interpretable good weights for another neural network.	

INVITED TALKS

ARC-AGI Without Pretraining	NVIDIA LLM Reasoning Team, <i>Sept 2025</i>
Generating Interpretable Networks Using Hypernetworks	Stanford Enigma Project, <i>May 2025</i>
ARC-AGI Without Pretraining	Gwangju Institute of Science and Technology, <i>Apr 2025</i>

AWARDS AND HONORS

ARC Prize 2025: <i>3rd place Paper Award, \$5000 prize.</i>	<i>Dec 2025</i>
International Physics Olympiad: <i>Silver Medal. 2nd in Canada.</i>	<i>July 2019</i>
International Physics Olympiad: <i>Honorable Mention. 5th in Canada.</i>	<i>July 2018</i>
Citadel Securities PhD Summit Poster Competition: <i>3rd place, \$5000 prize.</i>	<i>March 2025</i>
MIT Battlecode swarm intelligence competition: <i>1st place, \$8000 prize.</i>	<i>Jan 2022</i>